Track Number	Speaker Name	Time	Topic	Description
1	Ayyanar Jeyakrishnan	10:40 - 11:10 IST	Orchestrating Intelligence: Context Engineering Strategies for Multi-Agent Workflow	This session delves into context engineering techniques that enable effective orchestration of multi-agent workflows. It covers how agents can share memory, align goals, and coordinate tasks using structured prompts and contextual cues. The strategies discussed help build scalable, intelligent systems for complex, collaborative environments across domains like finance, healthcare, and logistics. What will be learned by end of session: Attendees will learn how to design and implement effective context engineering strategies to orchestrate multi-agent systems. They will understand how to think beyond prompt engineering or templates and enable collaboration through shared memory, structured prompts, and contextual alignment—empowering scalable, intelligent solutions for real-world, multi-domain challenges
1	Charan Panthangi	11:15 - 11:45 IST	From Prompting to Context Engineering: How MCP is Powering the Next Wave of Agentic AI.	The shift from prompt engineering to context engineering marks a new era in AI. Large Language Models alone are not enough — they need structured, dynamic context to reason, plan, and act reliably. This talk explores how Context Engineering, combined with the Model Context Protocol (MCP), enables scalable agentic AI systems that seamlessly integrate tools, memory, and enterprise data. Attendees will learn how MCP standardizes connections between LLMs and external environments, reducing hallucinations, improving task orchestration, and unlocking autonomous, production-ready AI ecosystems.
1	Yogesh Kumar, Navaneethan Gopal	11:50 - 12:20 IST	AI Agent Lifecycle – Evaluation to Execution	Comprehensive understanding of AgentOps—its principles, lifecycle, evaluation strategies, and best practices for deploying and managing autonomous AI agents in real-world environments.
1	Abishek Subramanian	12:25 - 12:55 IST	AI Agents Unleashed: Build, Deploy, and Scale with Databricks	Explore the Databricks Mosaic AI Agent Framework, a comprehensive platform designed for building production-ready AI agents at enterprise scale. This session covers the framework's core capabilities including multi-step agent orchestration, state management, integration with Unity Catalog for secure data and function access, and MLflow-based lifecycle management for versioning and deployment. Learn how to leverage Vector Search for real-time data retrieval, AI Gateway to integrate any large language models or custom tools, and Model Serving for scalable agent deployment with endpoint autoscaling and streaming responses. Discover advanced techniques in agent evaluation using automated AI judges and human feedback to ensure accuracy, safety, and compliance via built-in governance and guardrails. This technical deep dive provides a blueprint to efficiently develop, monitor, and optimize AI agents for robust, data-driven applications.
1	Preethi Natarajan	14:30 - 15:00 IST	From Fear to Empowerment: Embracing AI as a Human-Centric Enabler	In a world increasingly shaped by artificial intelligence, many professionals grapple with uncertainty—fearing job displacement, skill redundancy, and the erosion of human values. This talk explores how we can shift the narrative from insecurity to empowerment by embracing AI as a tool that amplifies human potential rather than replacing it. Through real-world examples, strategic insights, and design principles, we'll uncover how AI can be built and deployed in ways that prioritize empathy, inclusivity, and collaboration. Attendees will gain practical strategies to navigate change, upskill with purpose, and lead with a human-first mindset in an AI-driven future.
1	Anandan Kumaran	15:05 - 15:35 IST	AI-Proof Yourself: Why Intention Is the Superpower Machines Can't Copy	"AI is faster, cheaper, and everywhere—but it can't want. It has no intention, no purpose. Without intention, intelligence is just automation." Gen-AI can automate, accelerate, and analyze—but it cannot intend. It has no purpose of its own. In a world where machines keep getting smarter, the real challenge isn't about keeping up with AI, it's about becoming AI-Proof. AI works on patterns from the past, but humans can design the future. Our ability to set intentions, imagine possibilities, and move with purpose is something no algorithm can replicate. This unique strength—intention-orientation—is the antidote to obsolescence. Instead of seeing AI as a competitor, what if you could make it your co-partner? By aligning human intention with machine intelligence, you stop fearing replacement and start designing futures where AI does the heavy lifting, while you direct the "why" and the "what next." Join this session to learn how to AI-Proof yourself with intention-driven thinking. Discover practical frameworks and real-world strategies that future-proof your role, amplify your creativity, and help you lead in the Gen-AI era. "This session will show you how to future-proof yourself by mastering intention as your unique edge—transforming AI from a competitor into your most powerful co-partner."
1	Dhaval Moliya	15:40 - 16:10 IST	AI Powered Agentic Dev Loop @ Atlassian Scale	Managing a massive monorepo with over 100 million lines of code presents significant infrastructure cost, scale, build time and productivity challenges across SDLC. In the inner dev-loop, this scale results in sluggish search functionality and delayed IntelliSense auto-complete suggestions within the IDE. Additionally, generating repository-level insights and executing deep code refactoring pose significant challenges. To tackle these challenges head-on, we built AI-powered coding assistant in-house, which enhances productivity in various areas, including feature development, testing, documentation, code review, and maintenance. This tool creates a multiplier effect, alleviating the cognitive load on developers as millions of lines of AI-generated code are seamlessly integrated and accepted by them. In outer dev-loop, our high volume of daily builds, each running ~500,000 tests, often led to redundant testing. By leveraging static code analysis and machine learning for predictive test selection, we now skip 50-95% of tests per change, dramatically optimizing build times. Further AI-powered tooling has reduced manual effort in code migrations by 50%. These AI/ML advancements in our developer tooling have yielded annual infrastructure cost savings of approximately $4 million and saved thousands of engineering hours.
1	Harun Raseed Basheer	16:15 - 16:45 IST	MCP Unleashed—Live Workflow Automation for Your Entire Tech Organization!	Experience automation like never before! See MCP in action—one Python app, live demo, instantly automating DevOps, SRE, Security, and QA workflows on Azure. From business email to deployment, watch your whole tech team work hands-free. No slides. Just real, actionable results. Don't miss the hands-on enterprise AI revolution!
2	Madan Sudarman	10:40 - 11:10 IST	FinOps for AI + AI for FinOps	Explore the synergy between FinOps and AI — how FinOps principles can optimize AI workloads and how AI can enhance FinOps practices through automation and intelligence. Learn strategies to manage the unique cost dynamics of AI and GenAI, from GPU utilization to token-based pricing. Discover how to balance innovation, efficiency, and fiscal accountability in the era of intelligent cloud operations.
2	Sriram Radhakrishnan	11:15 - 11:45 IST	FinOps for Everyone: Taming Cloud Spend & Driving Innovation	FinOps is a transformative cultural practice and operational framework for managing cloud costs and maximising business value, fostering financial accountability in variable spend models. It functions through three iterative phases: Inform (gaining visibility), Optimize (driving efficiency), and Operate (embedding practices). The discipline promotes collaboration and shared responsibility among engineering, finance, and business teams, with core principles centred on data-driven decisions and value realisation. Key practices involve setting up proactive alerts, effective budgeting, and accurate forecasting, alongside robust tagging for granular cost allocation. Optimization strategies target various cloud services, including compute, storage, network, containers, serverless, and AI, utilising methods such as rightsizing, commitment models (e.g., Reserved Instances, Savings Plans), and automation via Policy as Code and serverless functions. FinOps also addresses common challenges like bill shock and the issue of untagged resources, leveraging both native cloud and third-party tools for comprehensive management. Ultimately, FinOps is an ongoing journey of continuous improvement and governance, crucial for achieving sustainable cloud financial mastery.
2	Kannan K	11:50 - 12:20 IST	Fake It Till You Make It: Emulating Cloud with LocalStack for Microservices	Ever wished you could run AWS/Azure in your own backyard (or laptop) and test your microservices without racking up a cloud bill? This talk shows you how to do exactly that using LocalStack, that lets you spin up cloud resources your development machine to rapidly build and debug cloud-native microservices entirely offline, with LocalStack acting as the cloud's stunt double. Expect a technical session with live demos – by the end, you'll be confidently "faking" cloud locally to accelerate your development.
2	Bipul Kumar, Neetu SN	12:25 - 12:55 IST	Building Scalable Systems with Event Driven Architecture	Event-Driven Architecture (EDA) is transforming how modern AI Driven systems handle scalability, resilience, and real-time responsiveness. In this session, we will explore the core principles of EDA, practical design patterns, learn how to implement event-driven systems using tools like Kafka, RabbitMQ, or cloud-native solutions, and discover strategies to overcome common challenges like event ordering, idempotency, and system observability.
2	Prateek Mandloi	14:30 - 15:00 IST	Reduce On-Call fatigue and Improve Incident response times with Atlassian AIOps	This session will include examples from Atlassian's real-world implementation of AIOps in a large-scale enterprise, demonstrating how AIOps improved on-call productivity within Atlassian by reducing alert noise by ~85%, creating over 400k+ groups, particularly useful for high-volume customers having 5k+ alerts/year. AIOps can reduce MTTR by suggesting relevant resources from Integration with other Atlassian products (Jira, Confluence, etc) and Integration with collaboration tools like Docs, Slack, etc., as well as past similar alerts and incidents. The talk will cover how to collaborate with your own OpsGuide (agentic teammate) powered by Atlassian's Rovo platform during the crucial incident investigation period to help recover faster and subsequently boost productivity by creating AutoPIRs (post-incident review). This session will inspire attendees to rethink their operational strategies and suggest real-world, practical ways in which others can learn to build an ecosystem where incident management is less stressful that reduces overall MTTR. Join us to explore how AIOps can elevate your organization's incident management processes. The talk will cover examples of tangible benefits of adopting Atlassian AIOps in their organizations. Key Takeaways 1. AIOps can reduce alert noise significantly, thereby reducing alert fatigue for on-call engineers experiencing high alert volumes. Further, AIOps can reduce troubleshooting time and improve overall incident response times. 2. The session will provide tangible examples of AIOps from Atlassian practices, inspiring attendees to revisit operational and agentic collaboration strategies. 3. Learn how to utilize data from multiple sources to build a robust context for incident investigation. Explore how integrations with various 3P systems and platforms can help in a seamless incident management workflow.
2	Sandeep Kumar Prakash	15:05 - 15:35 IST	Serverless is Easy, Until It Isn't	Serverless makes everything look easy, fast prototypes, instant scale, no servers to manage. But simplicity can be deceptive. As systems evolve, the very qualities that made serverless appealing begin to introduce friction: invisible complexity, unpredictable behaviour, and architectural blind spots. This talk explores the hidden challenges that surface when serverless architectures mature, the patterns that quietly erode reliability, cost-efficiency, and developer clarity. Serverless is easy when you start. This session is about what happens after.
2	Saravanan Paramasivan	15:40 - 16:10 IST	Building a Cloud Native IaC Platform for an Enterprise - Multicloud strategy	As enterprises evolve to meet dynamic business demands, the ability to operate seamlessly across multiple cloud environments has become a strategic imperative. And enterprises are increasingly adopting multi-cloud strategies to unlock innovation, ensure resilience and accelerate transformation. This session explores how NIQ empowers its developers to seamlessly operate across Azure, Google Cloud, AWS and private cloud environments—without compromising on governance, security or performance.
2	Muhammed Riyaz	16:15 - 16:45 IST	AI-Driven Automation in Data Pipelines	Exploring how AI transforms ETL pipelines from rigid, manual systems into intelligent, self-optimising workflows using tools like Airflow, Spark, and AWS Glue.
3	Jones Zachariah Noel	10:40 - 11:10 IST	Beyond the Black Box: Gaining Observability into AI-Powered Applications with Amazon Bedrock	Amazon Bedrock and Amazon Nova model enables you to build AI powered applications and when they are running in production, this can encounter errors or even simply a better visibility into how the model is performing is important. In this session, we will look at why monitoring is the key for AI applications and when using Bedrock with applications, how AI monitoring gives you insights not only about errors but also AI model performance.
3	Rabi Sahu, Keshav Jain	11:15 - 11:45 IST	Demystifying GenAI Security: Aligning Governance, Compliance & Risk	The surge in generative AI is reshaping how organizations approach innovation, with enterprise adoption accelerating at an unprecedented pace. As more businesses integrate AI into core operations, they encounter a growing array of security and regulatory challenges—particularly around sensitive data handling, unauthorized AI usage ("shadow AI"), and the tightening grip of global compliance mandates. Security experts are increasingly concerned about the risk of data exposure through AI platforms, while many decision-makers acknowledge gaps in their understanding of regulatory frameworks governing AI. These dynamics underscore the need for a proactive, structured approach to AI governance. Key Pillars for a Secure AI Ecosystem: Foundational Readiness: Begin by organizing and classifying data, establishing clear access controls, and implementing identity governance protocols to minimize exposure. Visibility into Risk: Map data flows, monitor usage patterns, and assess vulnerabilities in AI systems to uncover threats before they escalate. Application & Data Defense: Apply real-time security controls against emerging AI-specific threats—including prompt manipulation and leakage of proprietary or personal information. Responsible Governance: Build agile governance frameworks capable of responding to policy changes, ethical dilemmas, and the misuse of AI-generated outputs. By weaving these principles into their AI strategy, organizations can harness innovation without compromising trust, transparency, or compliance.
3	Dr.Anusuya Kirubakaran	11:50 - 12:20 IST	Agents in Action: Shielding Operational Risk and Mitigating Losses in Banking	She explores how AI agents are transforming risk management in banking. Discover how intelligent automation can identify, assign, and manage controls to reduce operational risks and financial losses proactively. Learn how next-gen AI systems are helping banks move from reactive problem-solving to predictive prevention, making operations safer, smarter, and more resilient.
3	Kamesh Sampath	12:25 - 12:55 IST	Ingest Anything, Unify Everything: Apache NiFi + LLMs + Apache Iceberg Universal Pipelines	Every organization struggles with the same data integration nightmare: dozens of data sources, each with different schemas, requiring months of manual engineering work to transform into analytics-ready formats. What if we could eliminate this bottleneck entirely using modern open-source tools and AI? In this hands-on session, we'll build intelligent data pipelines that automatically transform inconsistent data uploads into unified, analytics-ready tables using Apache NiFi orchestration, LLM-powered schema inference, and Apache Iceberg's evolution capabilities. What You'll Learn: How to build self-adapting data pipelines that understand semantic meaning across different schemas Practical implementation of LLM integration for automated schema mapping and validation Apache Iceberg's schema evolution features for maintaining backward compatibility while adapting to new data structures Real-world architecture patterns for scaling intelligent data pipelines across enterprise environments Ready to transform months of manual data engineering into minutes of intelligent automation? Join me to build the future of data integration.
3	Krishnakumar Narayanan	14:30 - 15:00 IST	From Stability to Speed: The New DNA of Tech Giants	Speed has become the currency of business. "Move fast and break things" is no longer a motto — it's the operating system. In a world where "show, don't tell" defines product reviews and performance metrics, tech giants have rewritten their DNA: shifting from long-term planning and predictable growth to relentless innovation, rapid iteration, and bold risk-taking. This talk explores the cultural disruptions and strategic pivots happening at breakneck speed — and what they mean for the next generation of engineers. As expectations rise and timelines shrink, future tech leaders must learn to thrive in environments where adaptability, urgency, and continuous learning aren't optional — they're survival traits.
3	Abdul Rasheed Feroz Khan	15:05 - 15:35 IST	Migration as a Profit Engine: Unlocking Partner Growth with Azure	In today's AI-powered landscape, partners stand at the forefront of a once-in-a-generation opportunity to shape the future of business. With Microsoft Azure, they're uniquely positioned to guide customers through transformative journeys—while unlocking new, sustainable revenue streams. This session will explore how cloud migration and modernization are no longer just technical initiatives—they are strategic growth engines. Discover how forward-thinking partners are leveraging Azure to drive innovation, deliver differentiated value, and lead the next wave of cloud adoption with confidence and profitability.
3	Sameer Mhaisekar	15:40 - 16:10 IST	Observing the outage: a case for resilient Observability	In this session we will discuss how to build observability strategy that is resilient and robust. We will discuss "earthquake signals" and "doom signals" that indicate imminent failure in systems and how we can make sure to keep an eye out and be prepared instead of being caught by surprise. This is a cloud-agnostic topic and covers best practices regardless of where the apps are running.
4	Jayakumar Vaithiyashankar	10:40 - 11:10 IST	From Code to Qubit: A Developer's Guide to Quantum Computing with IBM Qiskit	Explore the world of quantum computing! This talk, 'Beyond the Bit,' introduces the basic principles and provides practical skills to use them with IBM's open-source framework, Qiskit. The talk starts with the core concepts of quantum mechanics: superposition and entanglement. It moves beyond the limitations of traditional computing. Through examples and demonstrations using Qiskit and the IBM Quantum Composer, you will learn to construct quantum circuits, explore multi-qubit states, and implement key quantum gates and algorithms, such as the Bernstein-Vazirani algorithm and quantum key distribution. Real-world applications where quantum computing could lead to breakthroughs, including cryptography, drug discovery, and artificial intelligence, are examined. Qiskit's use in these fields is explored. This session is for developers, researchers, and anyone wanting to explore the frontier of computation and discover the potential that lies beyond the bit. You will gain a foundation in quantum programming with Qiskit and be inspired to explore this rapidly evolving field further
4	Karthikeyan VK	11:15 - 11:45 IST	Fine-Tuning LLMs: Build Your Own Domain Model	Session: Build Your Own GPT in 25 Minutes What if you could have a model that speaks your domain's language — whether that's finance, healthcare, legal, or even your own project data? In this 25-minute session, we'll demystify the process of creating your own GPT. Starting from the basics, you'll see how a base model (like GPT-2) can be adapted into a domain-specialized LLM. This isn't about generic AI — it's about models that understand your context. What You'll Learn The three paths to fine-tuning: full fine-tuning, parameter-efficient fine-tuning (LoRA, QLoRA), and instruction tuning with RLHF. Trade-offs explained: when to go heavy, when to stay efficient. Live walkthrough: how to structure data and fine-tune a GPT for your domain. Blueprint takeaway: a simple, repeatable process to create your own domain GPT. Who Should Attend Developers and ML engineers curious about custom LLMs. Researchers and tech leaders exploring domain-specific AI. Startup founders and product builders seeking differentiation with tailored AI. In just 25 minutes, you'll walk away with clarity, confidence, and a clear path to building your own GPT.
4	Sathish Nadarajan	11:50 - 12:20 IST	Harness the Power of Generative AI in Pharma: Accelerate Drug Discovery and Development	This session will provide an in-depth exploration of the drug discovery and development process, highlighting its inherent complexities and the significant barriers that often delay the introduction of new therapies to the market. Despite advances in science, developing a new drug remains a time-consuming, resource-intensive, and high-risk endeavor for pharmaceutical companies. In this presentation, we will examine not only the traditional workflow—spanning target identification, molecule design, preclinical validation, clinical trials, and patient safety monitoring—but also the persistent challenges that limit speed and efficiency at each stage. Amidst these challenges, Generative AI has emerged as a transformative technology with the capability to accelerate and optimize nearly every phase of the drug discovery and development lifecycle. We will provide a comprehensive overview of where and how AI models are currently being deployed, from rapidly identifying promising biological targets, designing novel compounds, predicting clinical success, and identifying safety signals, to drastically reducing time-to-market for new therapeutics. Drawing on real-world examples and case studies, the session will showcase practical applications and future potential of AI-powered approaches. Participants will gain insights into how pharma companies worldwide are leveraging generative AI in various pillars of R&D—improving efficiency, reducing costs, and ultimately bringing life-changing treatments to patients more swiftly and safely. Attendees will leave with a clear understanding of the latest advancements in AI for drug discovery, tangible use cases across the entire development journey, and strategies to unlock the full potential of AI in their own organizations.
4	Marshall Nicholas	12:25 - 12:55 IST	Simple Question, Perfect Answer	Imagine asking an AI chatbot a simple question about your business data and getting the perfect answer. Sounds easy, right? In reality, natural language querying of databases is far more complex than it appears. Think about the question "Who is our top customer?" Do we mean top by revenue, by growth, by number of purchases, or something else entirely? The AI won't know unless we provide context. In our session, we will be demoing a robust Text2SQL Agentic Workflow in AWS Stack that can be customized for your business. In the session , we will cover key ideas on metadata management, providing LLMs business context and observability best practices. By the end, you'll leave equipped with practical strategies and inspiration for creating your scalable, business-context-aware text-to-SQL solution. Session Outline 1. Introduction: The Challenge of Structured Data Querying 2. LLM Limitations with Structured Datasets 3. SQL Generator LLM → Synthesizer Technique 4. Why Agents Are Overkill; Autonomy Slider Concept 5. Embeddings for RAG and Metric Definitions 6. Business Rules vs. Industry Rules; Proper Nouns Handling 7. Latency vs. Accuracy Tradeoff 8. Quality Metadata Management 9. Observability-Driven Iterative Improvement 10. Deep Dive: Tech Stack (AWS Bedrock + Llamaindex for Orchestration) 11. Copilot-Like Partial Autonomy Demo 12. Building Custom-Tuned LLMs: When and Why 13. Case Study & Live Demo 14. Q&A Target Audience: Anyone who wants to setup a Natural Language Querying Platform on their Business Data that predominantly resided in Relational Databases
4	Dr. Santosh Karthikeyan Viswanathan	14:30 - 15:00 IST	AI in Healthcare	The transformative potential of AI in Healthcare, use cases and Enterprise AI Governance for developing and deploying AI solutions in a regulatory environment.
4	Vikas Gautam, Shweta Kamath	15:05 - 15:35 IST	Ideation to Production with Azure AI Foundry	-- The session will start with introducing an ideal path to production --> Choosing LLM (choices, design considerations) --> Choosing Agentic framework (choose between Open source frameworks like LangGraph, Autogen, etc.) --> Choosing architecture (Single/Multi-Agent) --> Building Agentic Application from scratch --> Evaluation Strategy --> Implement Guardrails --> Scaling Agents --> AgentOps -- Build a banking assistant with Semantic Kernel --> understanding design choices like LLM, architecture and frameworks --> Code walkthrough and demo -- Next, take a deep dive into Evaluation techniques in Agentic AI --> LLMs as a Judge --> Human-in-the-loop -- Next, discuss a few Evaluation frameworks --> Azure AI Evaluation SDK --> Mosaic AI Evaluation --> Evaluating with Azure Foundry --> Code review and demo -- Next, implement Guardrails with AI Foundry -- Lastly, how do you easily scale out your application with DAPR -- Code walkthrough and demo
4	Natarajan Ganapathi	15:40 - 16:10 IST	AI Agents: Build, Orchestrate and Scale with Microsoft Agent Framework & Azure AI Foundry	Learn hands-on how Microsoft Agent Framework provides the core engine and tools for developing agent logic and orchestration, while Azure AI Foundry Agent Service offers the managed platform for deploying and scaling these agents with enterprise-grade features
4	Muthukumar Ganesan	16:15 - 16:45 IST	Advancing Nuclear Safety with AI: Predicting Lifespan of Critical Components	In Fast Breeder Reactors (FBRs), the primary sodium centrifugal pump plays a vital role in circulating liquid sodium for effective heat transfer. A failure in this component due to wear, corrosion, or shaft misalignment, can pose severe operational and safety risks. Historical nuclear incidents such as the Three Mile Island accident, where loss of coolant occurred due to pump failure, Crystal River, where a coolant pump shaft broke after abnormal vibration, and Grand Gulf, which experienced recurring failures of the emergency residual heat removal pump, all underscore the critical role of reliable pump health monitoring in nuclear safety. Over time, internal components such as bearings degrade due to excessive vibration, overheating, or chemical corrosion, which can be identified in advance through sensor signals like vibration, temperature, flow rate, and pressure. This talk focuses on leveraging machine learning approaches for Remaining Useful Life (RUL) estimation of centrifugal pumps using multivariate time series sensor data. A hybrid methodology involving physical simulation models is employed to generate degradation datasets representing real-world operating conditions. The session explains how this time series data is used to train deep learning models to learn degradation patterns and accurately predict the RUL, thereby enabling early failure prediction. The talk also covers deployment strategies, focusing on redundant and diversified deployment techniques to ensure system reliability in safety-critical environments. This approach ensures timely alerts, optimized maintenance schedules, and improved reactor safety. Outline of Talk: 1. Introduction to Fast Breeder Reactors (3 mins) a. Overview of three stage nuclear program b. Role of fast breeder reactors and its challenges 2. Centrifugal Pump: Role and Failure Modes (3 mins) a. Importance of sodium pump in reactor operation b. Common failure modes: misalignment, cavitation, bearing wear, corrosion, overheat...